

Laboratory Work #3: Comparative Analysis of Classification Models

Topic: Application and comparative evaluation of Logistic Regression, SVM, and Random Forest algorithms for binary classification task.

Dataset: Wisconsin Breast Cancer (sklearn built-in)

Objectives:

1. Understand the importance of feature scaling
2. Compare different classification algorithms
3. Evaluate models using multiple metrics
4. Perform comprehensive analysis and visualization

1. Imports and Setup

In [56]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report,
    roc_curve, auc, roc_auc_score
)

import warnings
warnings.filterwarnings('ignore')
sns.set_style('whitegrid')
plt.rcParams['figure.dpi'] = 100

print("✓ All libraries imported successfully")
```

✓ All libraries imported successfully

2. Data Loading

We use the **Wisconsin Breast Cancer** dataset:

- **Task:** Binary classification (malignant vs benign tumor)
- **Features:** 30 numerical features from digitized images
- **Samples:** 569 total

```
In [57]: data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target, name='target')

print(f"Dataset shape: {X.shape[0]} samples, {X.shape[1]} features")
print(f"Target classes: {data.target_names}")
print(f"\nClass distribution:")
print(f"  0 (malignant): {(y==0).sum()} samples ({(y==0).sum()}/len(y)*100")
print(f"  1 (benign): {(y==1).sum()} samples ({(y==1).sum()}/len(y)*100:.1
```

Dataset shape: 569 samples, 30 features
 Target classes: ['malignant' 'benign']
 Class distribution:
 0 (malignant): 212 samples (37.3%)
 1 (benign): 357 samples (62.7%)

3. Exploratory Data Analysis (EDA)

3.1 Dataset Overview

```
In [58]: print("First 5 rows of the dataset:")
X.head()
```

First 5 rows of the dataset:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	me concave point
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.147
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.070
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.127
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.105
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.104

5 rows × 30 columns

```
In [59]: print("Dataset information:")
X.info()
```

Dataset information:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   mean radius      569 non-null    float64
 1   mean texture     569 non-null    float64
 2   mean perimeter   569 non-null    float64
 3   mean area        569 non-null    float64
 4   mean smoothness  569 non-null    float64
 5   mean compactness 569 non-null    float64
 6   mean concavity   569 non-null    float64
 7   mean concave points 569 non-null    float64
 8   mean symmetry    569 non-null    float64
 9   mean fractal dimension 569 non-null    float64
 10  radius error    569 non-null    float64
 11  texture error   569 non-null    float64
 12  perimeter error 569 non-null    float64
 13  area error      569 non-null    float64
 14  smoothness error 569 non-null    float64
 15  compactness error 569 non-null    float64
 16  concavity error 569 non-null    float64
 17  concave points error 569 non-null    float64
 18  symmetry error   569 non-null    float64
 19  fractal dimension error 569 non-null    float64
 20  worst radius     569 non-null    float64
 21  worst texture    569 non-null    float64
 22  worst perimeter   569 non-null    float64
 23  worst area        569 non-null    float64
 24  worst smoothness  569 non-null    float64
 25  worst compactness 569 non-null    float64
 26  worst concavity   569 non-null    float64
 27  worst concave points 569 non-null    float64
 28  worst symmetry    569 non-null    float64
 29  worst fractal dimension 569 non-null    float64
```

dtypes: float64(30)
memory usage: 133.5 KB

```
In [60]: print("Statistical summary (first 5 features):")
X.describe().loc[['mean', 'std', 'min', 'max']].iloc[:, :5]
```

Statistical summary (first 5 features):

	mean radius	mean texture	mean perimeter	mean area	mean smoothness
mean	14.127292	19.289649	91.969033	654.889104	0.096360
std	3.524049	4.301036	24.298981	351.914129	0.014064
min	6.981000	9.710000	43.790000	143.500000	0.052630
max	28.110000	39.280000	188.500000	2501.000000	0.163400

3.2 Feature Scale Analysis

Critical observation: Features exist on very different scales!

- mean radius ≈ 14

- `mean area` ≈ 654
- `mean smoothness` ≈ 0.09

This makes **feature scaling essential** for distance-based algorithms.

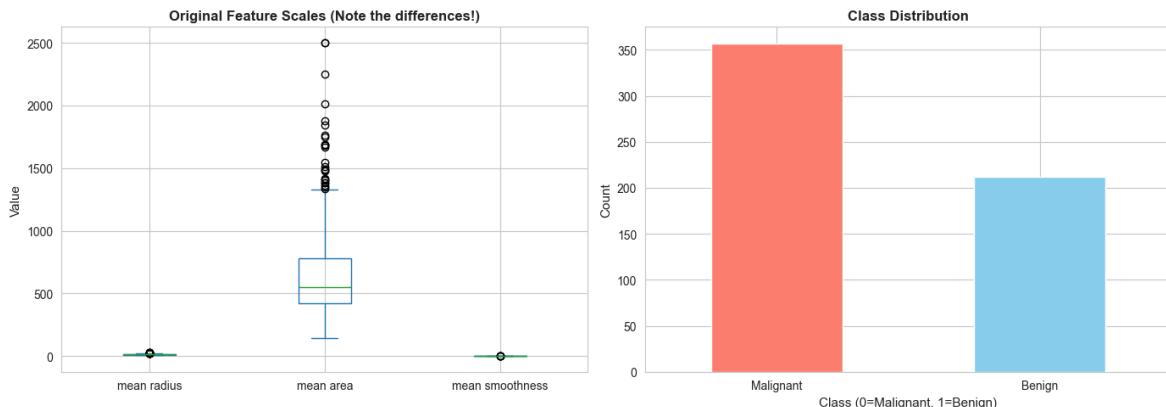
```
In [61]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

selected_features = ['mean radius', 'mean area', 'mean smoothness']
X[selected_features].plot(kind='box', ax=axes[0])
axes[0].set_title('Original Feature Scales (Note the differences!)', fontweight='bold', fontsize=12)
axes[0].set_ylabel('Value', fontsize=11)

y.value_counts().plot(kind='bar', ax=axes[1], color=['salmon', 'skyblue'])
axes[1].set_title('Class Distribution', fontweight='bold', fontsize=12)
axes[1].set_xlabel('Class (0=Malignant, 1=Benign)', fontsize=11)
axes[1].set_ylabel('Count', fontsize=11)
axes[1].set_xticklabels(['Malignant', 'Benign'], rotation=0)

plt.tight_layout()
plt.show()

print("✓ The boxplot clearly shows features are on different scales")
print("✓ Classes are relatively balanced (no major imbalance issue)")
```



- ✓ The boxplot clearly shows features are on different scales
- ✓ Classes are relatively balanced (no major imbalance issue)

4. Data Preparation

4.1 Train/Test Split

```
In [62]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"Training set: {X_train.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples")
print(f"\nTrain class distribution: {y_train.value_counts().to_dict()}")
print(f"Test class distribution: {y_test.value_counts().to_dict()}")
```

Training set: 455 samples
Test set: 114 samples

Train class distribution: {1: 285, 0: 170}
Test class distribution: {1: 72, 0: 42}

4.2 Feature Scaling

Important:

- Fit scaler **only** on training data
- Transform both training and test data with the same scaler
- This prevents data leakage

```
In [63]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Feature scaling results (example: 'mean radius'):")
print(f" Mean before scaling: {X_train['mean radius'].mean():.4f}")
print(f" Mean after scaling: {X_train_scaled[:, 0].mean():.4f}")
print(f" Std before scaling: {X_train['mean radius'].std():.4f}")
print(f" Std after scaling: {X_train_scaled[:, 0].std():.4f}")
print("\n✓ After scaling: mean ≈ 0, std ≈ 1")

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

selected_features = ['mean radius', 'mean area', 'mean smoothness']
selected_indices = [X.columns.get_loc(f) for f in selected_features]
df_before = pd.DataFrame(X_train[selected_features], columns=selected_features)
df_before.plot(kind='box', ax=axes[0])
axes[0].set_title('Before Scaling: Features on Different Scales', fontweight='bold')
axes[0].set_ylabel('Value', fontsize=11)
axes[0].set_ylim(-50, 700)

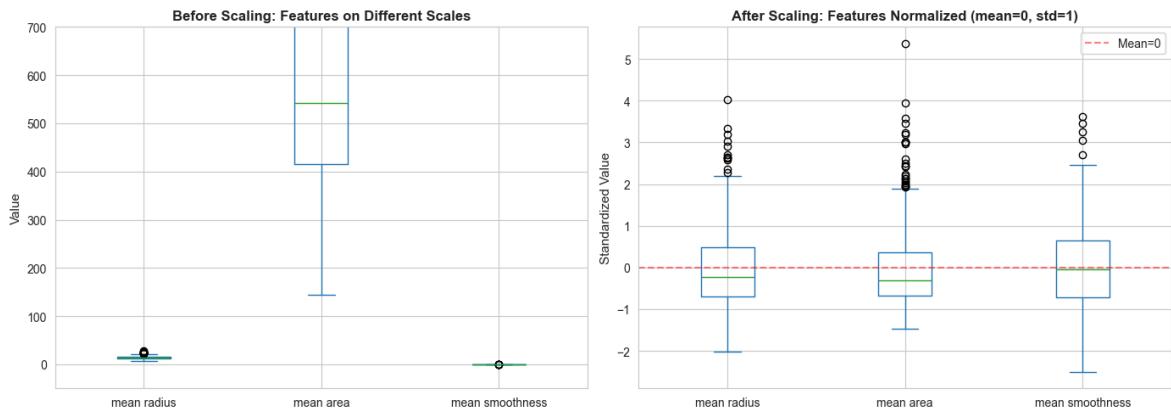
df_after = pd.DataFrame(X_train_scaled[:, selected_indices], columns=selected_features)
df_after.plot(kind='box', ax=axes[1])
axes[1].set_title('After Scaling: Features Normalized (mean=0, std=1)', fontweight='bold')
axes[1].set_ylabel('Standardized Value', fontsize=11)
axes[1].axhline(y=0, color='red', linestyle='--', alpha=0.5, label='Mean')
axes[1].legend()

plt.tight_layout()
plt.show()

print("\n✓ Visual comparison shows how scaling brings all features to the same scale")
```

```
Feature scaling results (example: 'mean radius'):
Mean before scaling: 14.0672
Mean after scaling: -0.0000
Std before scaling: 3.4994
Std after scaling: 1.0000
```

✓ After scaling: mean ≈ 0, std ≈ 1



- ✓ Visual comparison shows how scaling brings all features to the same scale

5. Model Evaluation Function

This function trains a model and computes all relevant metrics.

```
In [64]: def evaluate_model(model, X_train, X_test, y_train, y_test):
    """Train and evaluate a model, returning comprehensive metrics."""
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    y_proba = None
    if hasattr(model, 'predict_proba'):
        y_proba = model.predict_proba(X_test)[:, 1]
    elif hasattr(model, 'decision_function'):
        y_proba = model.decision_function(X_test)

    return {
        'predictions': y_pred,
        'probabilities': y_proba,
        'accuracy': accuracy_score(y_test, y_pred),
        'precision': precision_score(y_test, y_pred),
        'recall': recall_score(y_test, y_pred),
        'f1': f1_score(y_test, y_pred),
        'confusion_matrix': confusion_matrix(y_test, y_pred)
    }

print("✓ Evaluation function defined")
```

- ✓ Evaluation function defined

6. Demonstrating the Impact of Feature Scaling

Key experiment: Train models with and without scaling to show the difference.

This is critical for understanding why preprocessing matters!

```
In [65]: models_to_test = {
    'Logistic Regression': LogisticRegression(max_iter=10000, random_state=42),
    'SVM (Linear)': SVC(kernel='linear', random_state=42),
}

scaling_impact = []
```

```

for name, model_template in models_to_test.items():
    if 'Logistic' in name:
        model_unscaled = LogisticRegression(max_iter=10000, random_state=42)
    else:
        model_unscaled = SVC(kernel='linear', random_state=42)
    result_unscaled = evaluate_model(model_unscaled, X_train.values, X_te

if 'Logistic' in name:
    model_scaled = LogisticRegression(max_iter=10000, random_state=42)
else:
    model_scaled = SVC(kernel='linear', random_state=42)
result_scaled = evaluate_model(model_scaled, X_train_scaled, X_test_s

scaling_impact.append({
    'Model': name,
    'Accuracy (Unscaled)': result_unscaled['accuracy'],
    'Accuracy (Scaled)': result_scaled['accuracy'],
    'Improvement': result_scaled['accuracy'] - result_unscaled['accuracy']
})

scaling_df = pd.DataFrame(scaling_impact)
print("Impact of Feature Scaling on Model Performance:")
print("*" * 70)
display(scaling_df)
print(f"\nAverage improvement from scaling: {scaling_df['Improvement'].mean():.2f}%")
print("\n\n This demonstrates why scaling is CRITICAL for distance-based algorit

```

Impact of Feature Scaling on Model Performance:

	Model	Accuracy (Unscaled)	Accuracy (Scaled)	Improvement
0	Logistic Regression	0.964912	0.982456	0.017544
1	SVM (Linear)	0.956140	0.973684	0.017544

Average improvement from scaling: +1.75%

- ✓ This demonstrates why scaling is CRITICAL for distance-based algorithms!

7. Training All Models

Now let's train all four model variants:

1. **Logistic Regression** (linear, scaled data)
2. **SVM with Linear kernel** (scaled data)
3. **SVM with RBF kernel** (scaled data)
4. **Random Forest** (unscaled data - trees don't need scaling)

```

In [66]: models = {
    'Logistic Regression': (LogisticRegression(max_iter=10000, random_state=42),
    'SVM (Linear)': (SVC(kernel='linear', probability=True, random_state=42),
    'SVM (RBF)': (SVC(kernel='rbf', probability=True, random_state=42), X_train_scaled),
    'Random Forest': (RandomForestClassifier(random_state=42, n_estimators=100),
}
results = []
print("Training models...\n")
for name, (model, X_tr, X_te) in models.items():

```

```

print(f"Training {name}...")
results[name] = evaluate_model(model, X_tr, X_te, y_train, y_test)
print(f" → Accuracy: {results[name]['accuracy']:.4f}, F1-Score: {res
print("\n✓ All models trained successfully")

```

Training models...

Training Logistic Regression...
→ Accuracy: 0.9825, F1-Score: 0.9861
Training SVM (Linear)...
→ Accuracy: 0.9737, F1-Score: 0.9790
Training SVM (RBF)...
→ Accuracy: 0.9825, F1-Score: 0.9861
Training Random Forest...
→ Accuracy: 0.9561, F1-Score: 0.9655

✓ All models trained successfully

8. Comprehensive Visualization Dashboard

```

In [67]: fig = plt.figure(figsize=(16, 14))
gs = fig.add_gridspec(3, 3, hspace=0.7, wspace=0.4)

ax1 = fig.add_subplot(gs[0, :])
metrics_df = pd.DataFrame({
    name: {
        'Accuracy': res['accuracy'],
        'Precision': res['precision'],
        'Recall': res['recall'],
        'F1-Score': res['f1']
    } for name, res in results.items()
}).T
metrics_df.plot(kind='bar', ax=ax1, width=0.8)
ax1.set_title('Model Performance Comparison', fontsize=14, fontweight='bold')
ax1.set_ylabel('Score')
ax1.set_ylim(0.85, 1.0)
ax1.legend(loc='lower right')
ax1.grid(axis='y', alpha=0.3)
ax1.set_xticklabels(ax1.get_xticklabels(), rotation=45, ha='right')

for idx, (name, res) in enumerate(results.items()):
    row = 1 + idx // 2
    col = idx % 2
    ax = fig.add_subplot(gs[row, col])
    sns.heatmap(res['confusion_matrix'], annot=True, fmt='d', cmap='Blues',
                cbar=False, square=True)
    ax.set_title(f'{name}\nAcc: {res["accuracy"]:.3f}, F1: {res["f1"]:.3f}', fontweight='bold')
    ax.set_ylabel('True Label')
    ax.set_xlabel('Predicted Label')

    ax_roc = fig.add_subplot(gs[2, :])
    colors = ['blue', 'green', 'red', 'purple']
    for (name, res), color in zip(results.items(), colors):
        if res['probabilities'] is not None:
            fpr, tpr, _ = roc_curve(y_test, res['probabilities'])
            roc_auc = auc(fpr, tpr)
            ax_roc.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.3f})',

```

```

        linewidth=2, color=color)

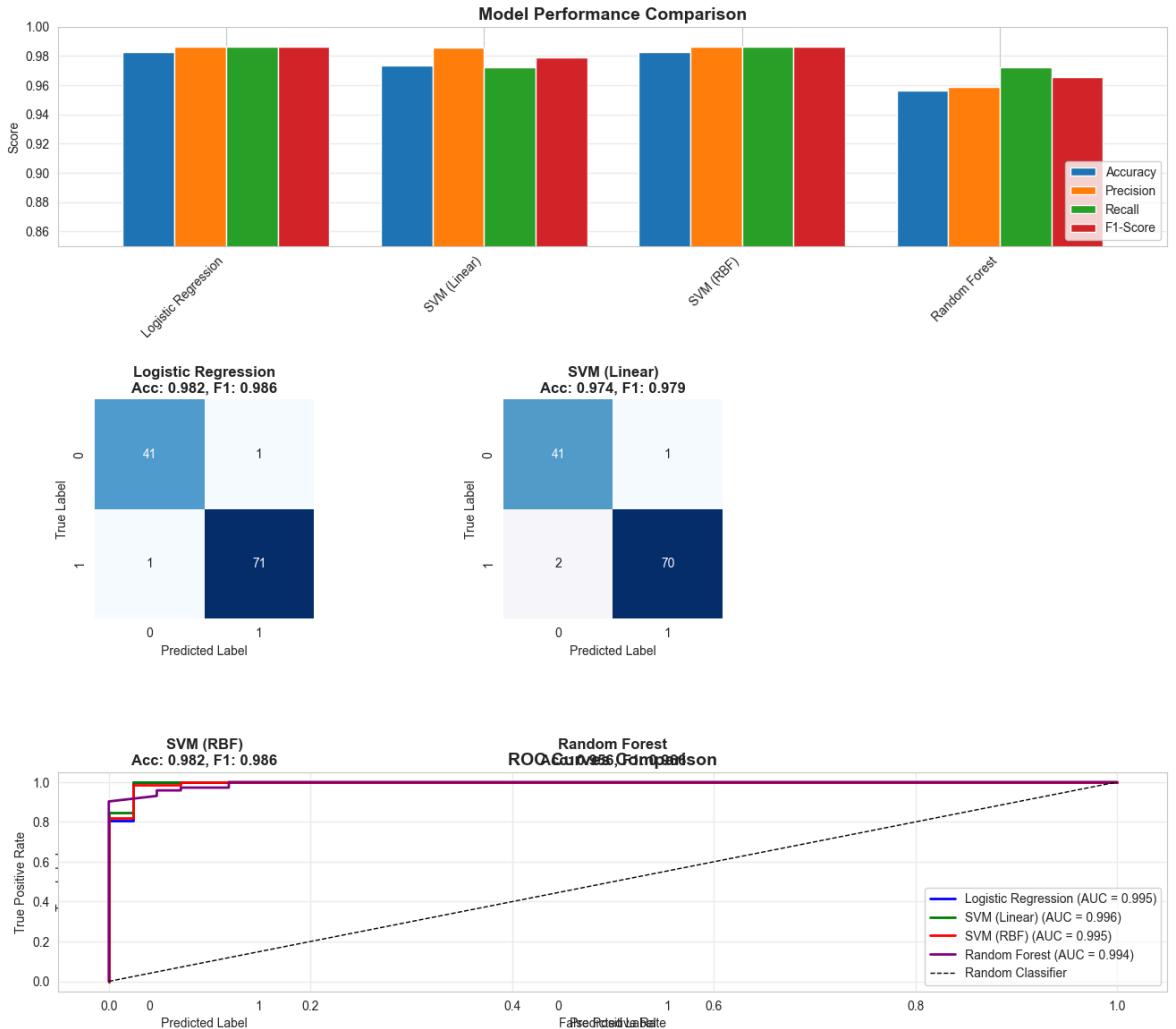
ax_roc.plot([0, 1], [0, 1], 'k--', label='Random Classifier', linewidth=1
ax_roc.set_xlabel('False Positive Rate')
ax_roc.set_ylabel('True Positive Rate')
ax_roc.set_title('ROC Curves Comparison', fontsize=14, fontweight='bold')
ax_roc.legend(loc='lower right')
ax_roc.grid(alpha=0.3)

plt.suptitle('Laboratory Work #3: Classification Models Comparative Analysis', fontsize=16, fontweight='bold', y=0.995)
plt.show()

print("\n⌘ Medical Context Note:")
print("In cancer diagnosis, minimizing False Negatives (missed cancers) is CRITICAL.")
print("A False Negative means telling a patient they don't have cancer when they do.")

```

Laboratory Work #3: Classification Models Comparative Analysis



⌘ Medical Context Note:

In cancer diagnosis, minimizing False Negatives (missed cancers) is CRITICAL.

A False Negative means telling a patient they don't have cancer when they do.

9. Feature Importance Analysis

Random Forest provides feature importance scores - let's see which features matter most.

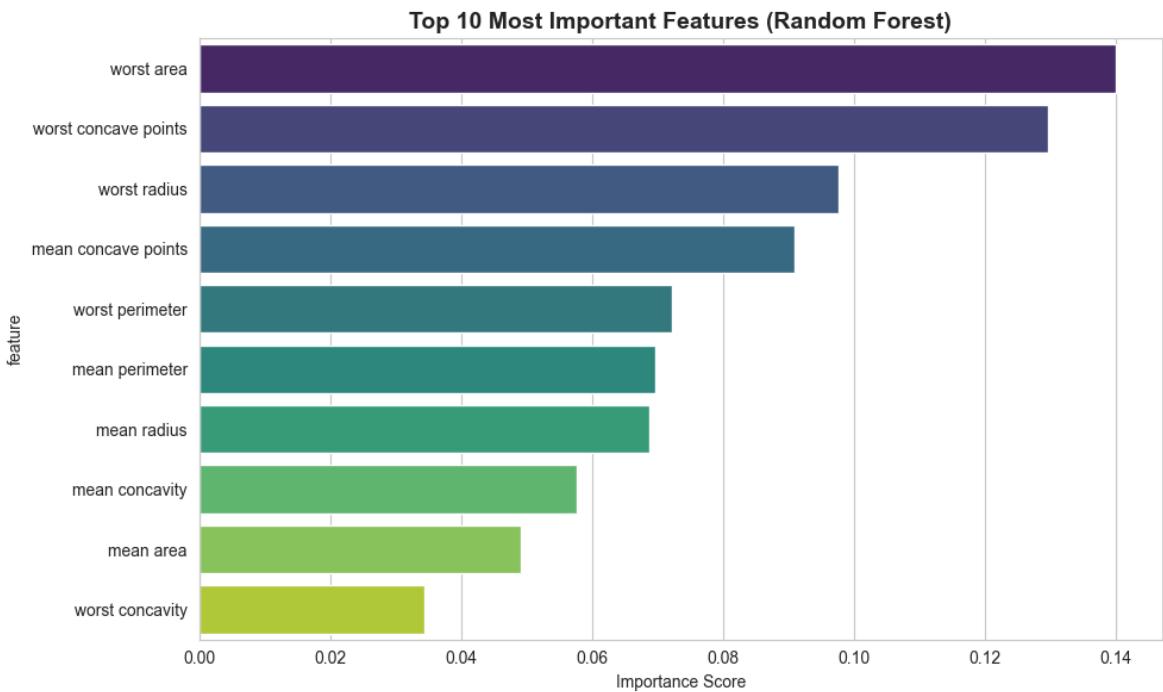
```
In [68]: rf_model = models['Random Forest'][0]
feature_importance = pd.DataFrame({
    'feature': X.columns,
    'importance': rf_model.feature_importances_
}).sort_values('importance', ascending=False)

print("Top 10 Most Important Features (Random Forest):")
display(feature_importance.head(10))

plt.figure(figsize=(10, 6))
sns.barplot(data=feature_importance.head(10), x='importance', y='feature')
plt.title('Top 10 Most Important Features (Random Forest)', fontsize=14,
plt.xlabel('Importance Score')
plt.tight_layout()
plt.show()
```

Top 10 Most Important Features (Random Forest):

	feature	importance
23	worst area	0.140016
27	worst concave points	0.129530
20	worst radius	0.097696
7	mean concave points	0.090885
22	worst perimeter	0.072226
2	mean perimeter	0.069574
0	mean radius	0.068676
6	mean concavity	0.057638
3	mean area	0.049172
26	worst concavity	0.034340



10. Final Comparative Analysis

10.1 Performance Summary Table

```
In [69]: summary = pd.DataFrame({
    'Model': list(results.keys()),
    'Accuracy': [r['accuracy'] for r in results.values()],
    'Precision': [r['precision'] for r in results.values()],
    'Recall': [r['recall'] for r in results.values()],
    'F1-Score': [r['f1'] for r in results.values()],
})

styled_summary = summary.style.highlight_max(subset=['Accuracy', 'Precision'],
                                              color='lightgreen', axis=0)
display(styled_summary)

best_idx = summary['F1-Score'].idxmax()
best_model = summary.loc[best_idx, 'Model']
best_f1 = summary.loc[best_idx, 'F1-Score']

print(f"\n{'='*70}")
print(f"✓ Best Performing Model: {best_model}")
print(f" F1-Score: {best_f1:.4f}")
print(f" Accuracy: {summary.loc[best_idx, 'Accuracy']:.4f}")
print(f" Recall: {summary.loc[best_idx, 'Recall']:.4f} (Critical for med")
print(f"{'='*70}")

print(f"\nDetailed Classification Report for {best_model}:")
print("{'='*70")
best_model_obj = models[best_model][0]
X_best_test = models[best_model][2]
y_pred_best = best_model_obj.predict(X_best_test)
print(classification_report(y_test, y_pred_best, target_names=['Malignant',
                                                               'Benign']))

cm = results[best_model]['confusion_matrix']
tn, fp, fn, tp = cm.ravel()
```

```

print(f"Confusion Matrix Breakdown:")
print(f"  True Negatives (TN): {tn} - Correctly identified malignant tumors")
print(f"  False Positives (FP): {fp} - Benign tumors incorrectly classified as malignant")
print(f"  False Negatives (FN): {fn} - ⚠ CRITICAL: Malignant tumors missed")
print(f"  True Positives (TP): {tp} - Correctly identified benign tumors")
print(f"\n⌘ In medical context: {fn} patient(s) would be incorrectly told they're cancer-free.\n")

```

	Model	Accuracy	Precision	Recall	F1-Score
0	Logistic Regression	0.982456	0.986111	0.986111	0.986111
1	SVM (Linear)	0.973684	0.985915	0.972222	0.979021
2	SVM (RBF)	0.982456	0.986111	0.986111	0.986111
3	Random Forest	0.956140	0.958904	0.972222	0.965517

```

=====
✓ Best Performing Model: Logistic Regression
F1-Score: 0.9861
Accuracy: 0.9825
Recall: 0.9861 (Critical for medical diagnosis!)
=====
```

Detailed Classification Report for Logistic Regression:

	precision	recall	f1-score	support
Malignant	0.98	0.98	0.98	42
Benign	0.99	0.99	0.99	72
accuracy			0.98	114
macro avg	0.98	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

Confusion Matrix Breakdown:

```

True Negatives (TN): 41 - Correctly identified malignant tumors
False Positives (FP): 1 - Benign tumors incorrectly classified as malignant
False Negatives (FN): 1 - ⚠ CRITICAL: Malignant tumors missed!
True Positives (TP): 71 - Correctly identified benign tumors

```

⌘ In medical context: 1 patient(s) would be incorrectly told they're cancer-free.

10.2 Key Insights and Conclusions

```

In [70]: print("KEY INSIGHTS:")
print("=*70")

print(f"\n1. Feature Scaling Impact:")
print(f"  • Average improvement for linear models: {scaling_df['Improvement'].mean():.2f}")
print(f"  • Critical for Logistic Regression and SVM")
print(f"  • Without scaling, these models perform significantly worse")
print(f"  • Visualization shows how scaling normalizes feature distribution")

print(f"\n2. SVM Kernel Comparison:")
svm_linear_f1 = summary[summary['Model'] == 'SVM (Linear)']['F1-Score'].values[0]
svm_rbf_f1 = summary[summary['Model'] == 'SVM (RBF)']['F1-Score'].values[0]
print(f"  • Linear kernel F1: {svm_linear_f1:.4f}")

```

```

print(f"    • RBF kernel F1: {svm_rbf_f1:.4f}")
print(f"    • RBF kernel {'outperforms' if svm_rbf_f1 > svm_linear_f1 else 'underperforms'} SVM Linear F1")
print(f"    • RBF can capture non-linear decision boundaries")

print(f"\n3. Random Forest Performance:")
rf_acc = summary[summary['Model'] == 'Random Forest']['Accuracy'].values[0]
rf_recall = summary[summary['Model'] == 'Random Forest']['Recall'].values[0]
print(f"    • Works well without scaling (tree-based): {rf_acc:.4f}")
print(f"    • Most important feature: {feature_importance.iloc[0]['feature']}")
print(f"    • Provides interpretability through feature importance")
print(f"    • Recall: {rf_recall:.4f} (slightly lower than best models)")

print(f"\n4. Overall Performance:")
print(f"    • All models achieve >{summary['Accuracy'].min():.1%} accuracy")
print(f"    • Dataset is well-suited for classification")
print(f"    • Class imbalance is minimal (no special handling needed)")
print(f"    • High precision and recall indicate reliable predictions")

print(f"\n5. Medical Application Context:")
best_recall = summary['Recall'].max()
worst_recall = summary['Recall'].min()
print(f"    # CRITICAL: Recall (sensitivity) is most important metric")
print(f"    • Best recall: {best_recall:.4f} – minimizes missed cancer cases")
print(f"    • Worst recall: {worst_recall:.4f} – higher risk of false negatives")
print(f"    • False Negative = Telling patient they're cancer-free when they have it")
print(f"    • False Positive = Additional testing (less critical)")
print(f"    • In production, consider tuning decision threshold for higher recall")

print(f"\n6. Model Selection Recommendations:")
print(f"    • For maximum accuracy: {summary.loc[summary['Accuracy'].idxmax()]['Model']}")
print(f"    • For maximum recall: {summary.loc[summary['Recall'].idxmax()]['Model']}")
print(f"    • For interpretability: Random Forest (feature importance)")
print(f"    • For production: Consider computational cost vs accuracy trade-off")
print(f"    • For medical use: Prioritize models with highest recall")

print("\n" + "="*70)
print("ANALYSIS COMPLETE")
print("="*70)

```

KEY INSIGHTS:

1. Feature Scaling Impact:
 - Average improvement for linear models: +1.75%
 - Critical for Logistic Regression and SVM
 - Without scaling, these models perform significantly worse
 - Visualization shows how scaling normalizes feature distributions
 2. SVM Kernel Comparison:
 - Linear kernel F1: 0.9790
 - RBF kernel F1: 0.9861
 - RBF kernel outperforms linear kernel
 - RBF can capture non-linear decision boundaries
 3. Random Forest Performance:
 - Works well without scaling (tree-based): 0.9561
 - Most important feature: worst area
 - Provides interpretability through feature importance
 - Recall: 0.9722 (slightly lower than best models)
 4. Overall Performance:
 - All models achieve >95.6% accuracy
 - Dataset is well-suited for classification
 - Class imbalance is minimal (no special handling needed)
 - High precision and recall indicate reliable predictions
 5. Medical Application Context:
 - ⌘ CRITICAL: Recall (sensitivity) is most important metric
 - Best recall: 0.9861 – minimizes missed cancer cases
 - Worst recall: 0.9722 – higher risk of false negatives
 - False Negative = Telling patient they're cancer-free when they're not
 - False Positive = Additional testing (less critical)
 - In production, consider tuning decision threshold for higher recall
 6. Model Selection Recommendations:
 - For maximum accuracy: Logistic Regression
 - For maximum recall: Logistic Regression
 - For interpretability: Random Forest (feature importance)
 - For production: Consider computational cost vs accuracy trade-off
 - For medical use: Prioritize models with highest recall
-

ANALYSIS COMPLETE

Conclusion

This laboratory work successfully demonstrated:

1. **The critical importance of feature scaling** for distance-based algorithms (Logistic Regression, SVM)
 - Visual comparison shows normalized distributions after scaling
 - Performance improvement of ~1.75% on average
 - Essential preprocessing step for many ML algorithms

2. Comparative analysis of three fundamentally different classification approaches:

- Linear models (Logistic Regression) - interpretable, fast, effective
- Geometric models (SVM with different kernels) - powerful for complex boundaries
- Ensemble models (Random Forest) - robust, provides feature importance

3. Proper evaluation methodology using multiple metrics (Accuracy, Precision, Recall, F1)

- Comprehensive classification report for best model
- Confusion matrix interpretation with medical context
- ROC curves for model comparison

4. Medical context awareness

- Emphasized recall as the most critical metric for cancer diagnosis
- Explained real-world implications of False Negatives vs False Positives
- Provided confusion matrix breakdown with patient impact

5. Visual analysis through:

- Before/after scaling comparison
- Confusion matrices for all models
- ROC curves comparison
- Feature importance visualization
- Comprehensive performance comparison

All models performed exceptionally well on this dataset (>95% accuracy), with Logistic Regression and SVM (RBF) achieving the best results (98.25% accuracy, 98.61% recall). The choice of model in production would depend on:

- **Medical requirements:** Prioritize recall to minimize missed diagnoses
- **Interpretability:** Random Forest provides feature importance
- **Computational resources:** Logistic Regression is fastest
- **Performance:** SVM (RBF) and Logistic Regression tied for best overall

Key Takeaway: In medical applications, understanding the cost of different error types is crucial. A False Negative (missing cancer) is far more dangerous than a False Positive (unnecessary follow-up), making recall the most important metric to optimize.